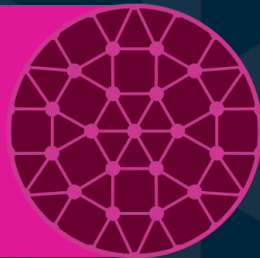


Tensorboard
with
run:ai
(using ResNet examples)

Tensorboard overview



What is needed for Tensorboard to run?

- 1) A **logs folder** to store objects related to the runs

Logs folder



How does Tensorboard work?

1) To **view the Tensorboard UI**, we start the server with a CLI command

You must specify the location of the logs folder with the logdir argument. You can also set the host IP and port.

```
tensorboard \  
  --logdir=/abs/path/to/logs_folder \  
  --host=0.0.0.0 \  
  --port=6006
```

How does Tensorboard work?

2) In order to **write records to the Tensorboard folder**, create a Tensorboard callback in our script, and pass it to a model

Note:

It is not necessary to start the Tensorboard server in order to write records.

Starting the server is only needed to view the UI.

```
import tensorflow as tf

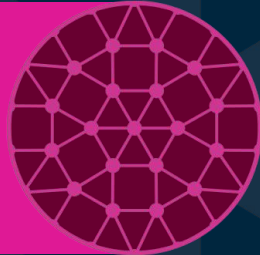
tb_dir = "/abs/path/to/logs_folder"

"""
code to build and compile your model
"""

tensorboard_callback = tf.keras.callbacks.TensorBoard(
    log_dir=tb_dir)

history = model.fit(train_ds,
                    epochs=5,
                    callbacks=[tensorboard_callback])
```

Tensorboard with run:ai



What is needed for Tensorboard to run on run:ai?

1) A **persistent** directory to keep

- Tensorboard logs folder

2) A **docker image** with the following installed

- Tensorboard
- jupyterlab*
- jupyter-server-proxy*


*needed to access the Tensorboard UI

3) A **docker image** with the following installed

- Tensorflow**
- Keras**

**needed in order train Tensorflow models (ResNet in our example)


Creating persistent directory



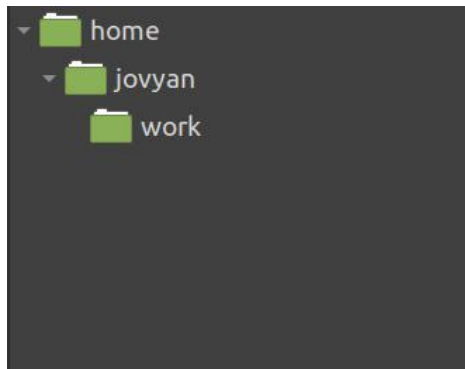
Before we start

We need to create a 'tensorboard_logs' folder on our NFS.

This is what our NFS folder structure looks like now



This is the default folder structure for our jupyter lab image

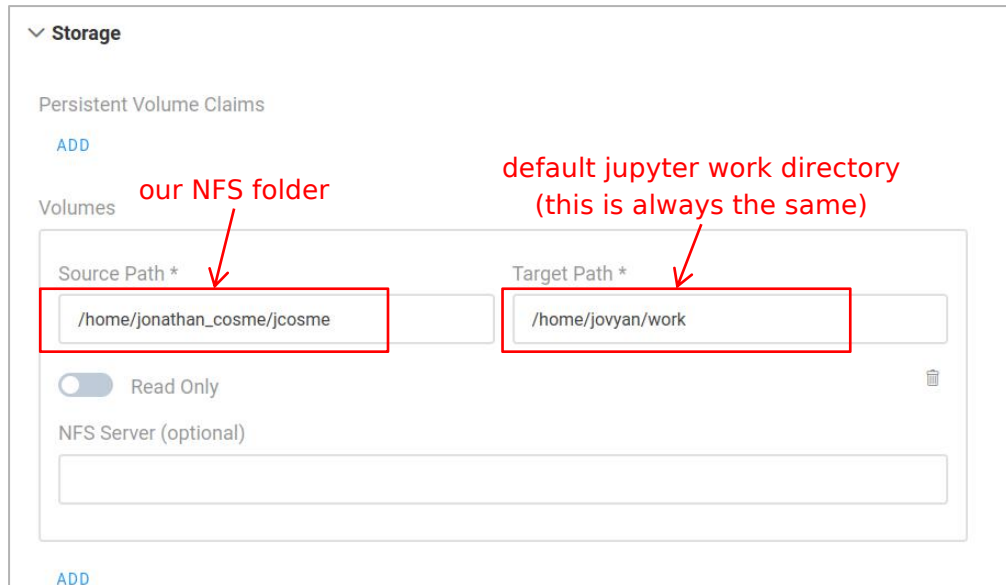


Before we start

We need to create an 'tensorboard_logs' folder on our NFS.

For *our example*:

Whenever we create a job on run:ai, we *must always* mount our NFS to the default jupyter work directory




The screenshot shows the 'Storage' configuration page in the run:ai interface. Under the 'Volumes' section, a new volume is being configured. The 'Source Path *' field is set to '/home/jonathan_cosme/jcosme', which is highlighted with a red box and labeled 'our NFS folder' with a red arrow. The 'Target Path *' field is set to '/home/jovyan/work', also highlighted with a red box and labeled 'default jupyter work directory (this is always the same)' with a red arrow. Below these fields, there is a 'Read Only' toggle switch (currently off) and an 'NFS Server (optional)' text input field. An 'ADD' button is visible at the bottom of the configuration area.


Before we start

We need to create an 'tensorboard_logs' folder on our NFS.


1. After we mount our NFS volume, our work directory will look like this:




2. Using Jupyter Lab, we create an 'tensorboard_logs' folder within the work directory:



3. This will cause our NFS directory to automatically look like this:



Docker images



First docker image used in our example

The docker image we will use to access the UI is:
jonathancosme/tensorboard-ui

This is what is in the dockerfile:

```
FROM jonathancosme/root-jpy-prox
RUN mamba install -c conda-forge tensorboard -y &&
    mamba clean --all -f -y
COPY jupyter_server_config.py /etc/jupyter/
```

← notebook with jupyter-server-proxy already installed

← install Tensorboard

← copy new config file, with Tensorboard UI access configured

First docker image used in our example

The docker image we will use to access the UI is:

jonathancosme/tensorboard-ui

In order to access the Tensorboard UI, we need to **add this entry to the jupyter_server_config.py** file, and replace the existing file in the image

```
c.ServerProxy.servers = {
    'tensorboard-server': {
        'command': [
            'tensorboard',
            '--logdir=/home/jovyan/work/tensorboard_logs',
            '--host=0.0.0.0',
            '--port=6006',
        ],
        'timeout': 30,
        'launcher_entry': {
            'title': 'tensorboard'
        },
    },
    'port': 6006
}
```

we specify our Tensorboard logs folder locations (this is why we must always mount our NFS directory to the default jupyter work directory)

we must make sure to start the server on this IP and port

This tells jupyter to forward port 6006 to the URL

Second docker image used in our example

The docker image we will use to train a Tensorflow ResNet model is:

jonathancosme/keras-nb


This is what is in the dockerfile:

```
FROM jonathancosme/base-notebook-root-py38
RUN wget https://developer.download.nvidia.com/compute/cuda/repos/ubuntu2004/x86_64/cuda-keyring_1.0-1_all.deb && \
    sudo DEBIAN_FRONTEND=noninteractive dpkg -i cuda-keyring_1.0-1_all.deb && \
    sudo apt-get update && \
    sudo DEBIAN_FRONTEND=noninteractive apt-get -y install cuda && \
    sudo apt-get autoclean
RUN conda config --set remote_read_timeout_secs 300 && \
    conda config --set remote_connect_timeout_secs 300 && \
    CONDA_OVERRIDE_CUDA="11.7" mamba install -c conda-forge -c nvidia numpy matplotlib pillow tensorflow-gpu=2.8 cudatoolkit=11.7 cudnn -y && \
    pip install tensorflow==2.9 tensorboard==2.9 tensorflow-hub tensorflow-text keras-nlp && \
    mamba clean --all -f -y
RUN export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/ && \
    export TF_FORCE_GPU_ALLOW_GROWTH=true
RUN sudo apt-get -y install libcudnn8-dev && \
    sudo apt-get autoclean && \
    sudo rm cuda-keyring_1.0-1_all.deb
```

Annotations:

- notebook with root privileges and python 3.8
- install cudatoolkit 11.7
- Install Tensorflow
- set environment variable needed for Tensorflow to run

Accessing the Tensorboard UI



Access Tensorboard UI

Create a jupyter interactive job with:

- image jonathancosme/tensorboard-ui
- mounted NFS folder (with 'tensorboard_logs' folder) in default jupyter work directory

INTERACTIVE TRAINING

Name Project * GPUs

testproj 0

Image *

jonathancosme/tensorboard-ui

Distributed Training (MPI) ☐

Jupyter Notebook ☒

Jupyter Notebook Password

Storage

Persistent Volume Claims

ADD

Volumes

Source Path * Target Path *

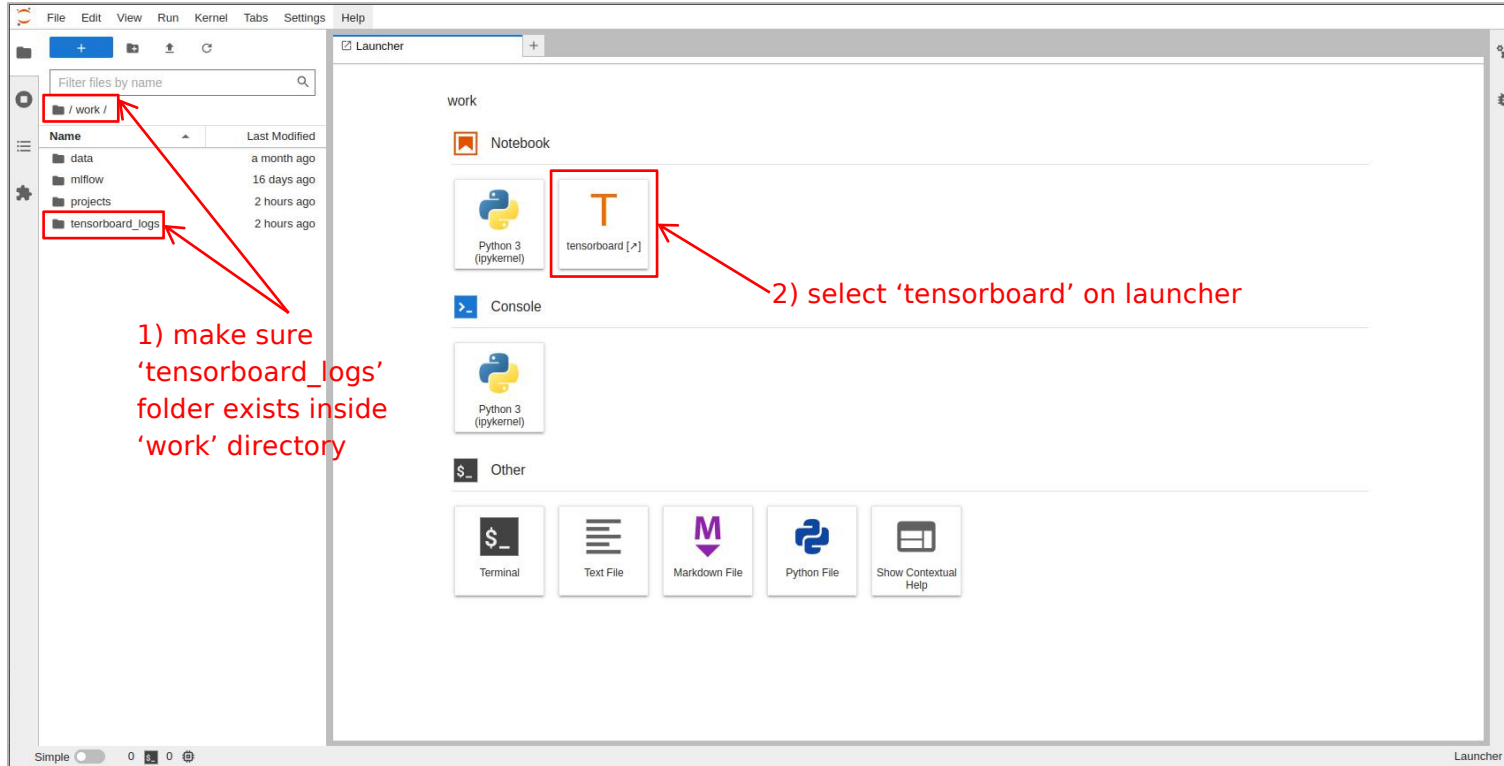
/home/jonathan_cosme/jcosme /home/jovyan/work

☒ Read Only

NFS Server (optional)

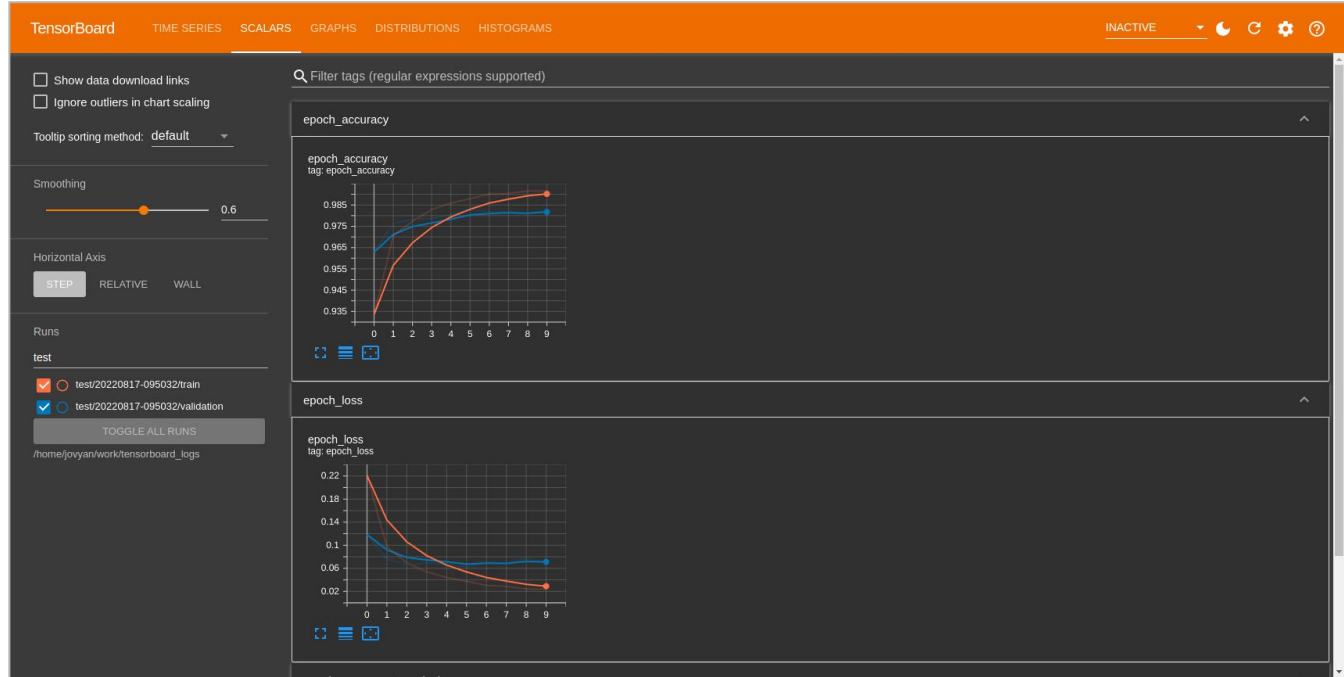
ADD

Access Tensorboard UI



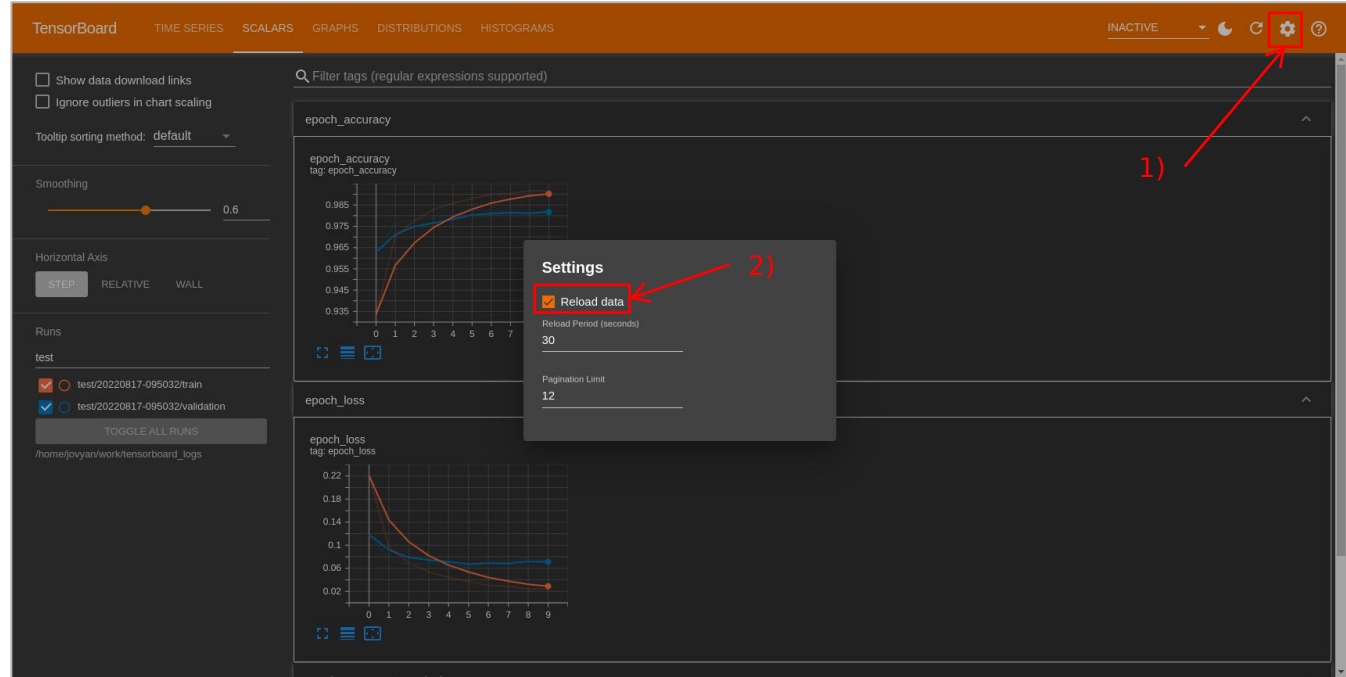
Access Tensorflow UI

A new tab should appear with the Tensorflow UI



Access Tensorflow UI

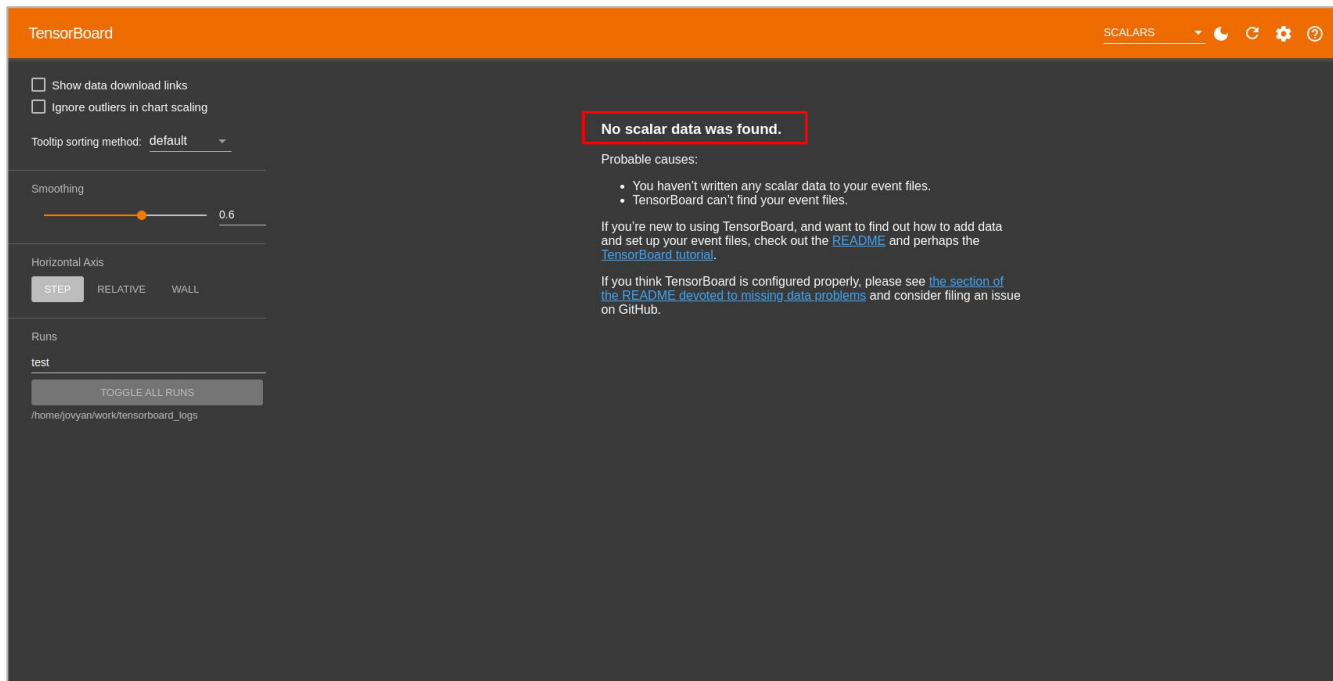
Select 'Reload data' under settings.




Access Tensorflow UI

Note:

The first time you access the UI, there will be no data available



running Tensorflow ResNet
experiments with run:ai



Python scripts (ResNet50 example)

Set the absolute path of the logs folder, and set a project name.

```
tb_dir = "/home/jovyan/work/tensorboard_logs"
project_name = 'resnet50'
```

Load the flowers dataset from Tensorflow

```
data_root = tf.keras.utils.get_file(
    'flower_photos',
    'https://storage.googleapis.com/download.tensorflow.org/example_images/flower_photos.tgz',
    cache_dir='.',
    untar=True)
```

Load the ResNet architecture directly from Tensorflow

```
model = tf.keras.applications.resnet.ResNet50(
    include_top=True,
    weights=None,
    input_shape=(224, 224, 3),
    classes=num_classes,
    classifier_activation='softmax',
    pooling='avg',
)
```

Create a subdirectory in our tensorboard_logs folder, using project_name, then create another subfolder using the date and time

```
log_dir = log_dir = f"{tb_dir}/{project_name}/{datetime.datetime.now().strftime('%Y%m%d-%H%M%S')}"
```

```
tensorboard_callback = tf.keras.callbacks.TensorBoard(
    log_dir=log_dir,
    histogram_freq=1)
```

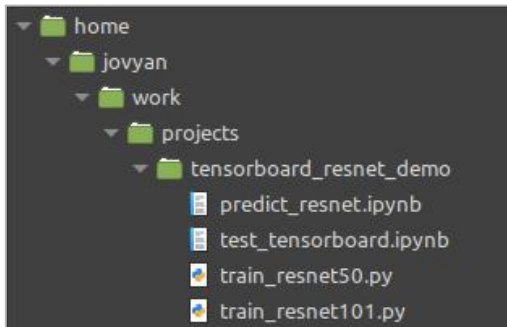
Pass in the callback, fit and save the model

```
history = model.fit(train_ds,
                    validation_data=val_ds,
                    epochs=NUM_EPOCHS,
                    callbacks=[tensorboard_callback])

model.save('./resnet50')
```

CLI submission

Our example scripts are located here:



so our CLI command would look like this:

```
runai submit \  
--project testproj \  
--gpu 1 \  
--job-name-prefix tb-resnet-demo \  
--image jonathancosme/keras-nb \  
--volume /home/jonathan_cosme/jcosme:/home/jovyan/work \  
-- python work/projects/tensorboard_resnet_demo/train_resnet50.py
```

make sure you use the keras-nb
docker image

make sure you mount the NFS to
the work directory



The code to run the job must specify
the python script with relative
location

Example job submission

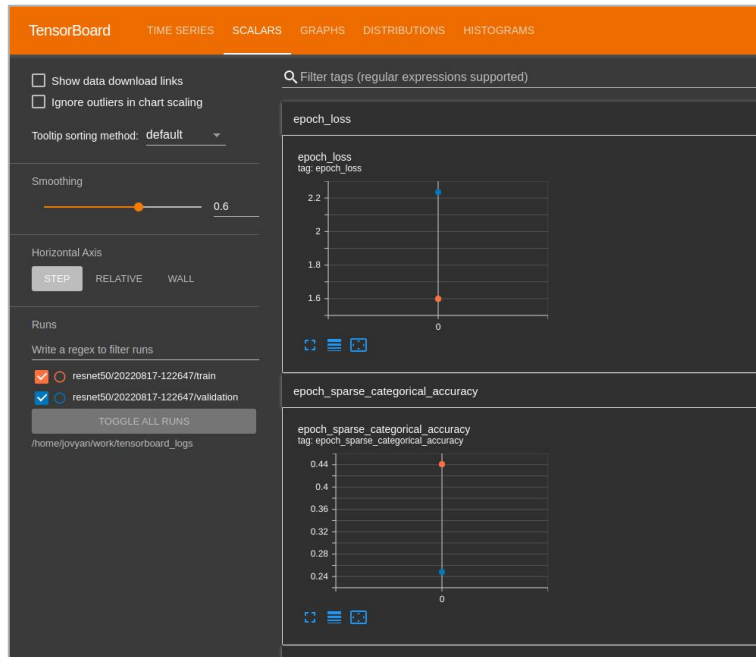
1) submit CLI command

```
(k8s) jcosme@jane:~$ runai submit \  
> --project testproj \  
> --gpu 1 \  
> --job-name-prefix tb-reset-demo \  
> --image jonathancosme/keras-nb \  
> --volume /home/jonathan_cosme/jcosme:/home/jovyan/work \  
> -- python work/projects/tensorboard_resnet_demo/train_resnet50.py
```

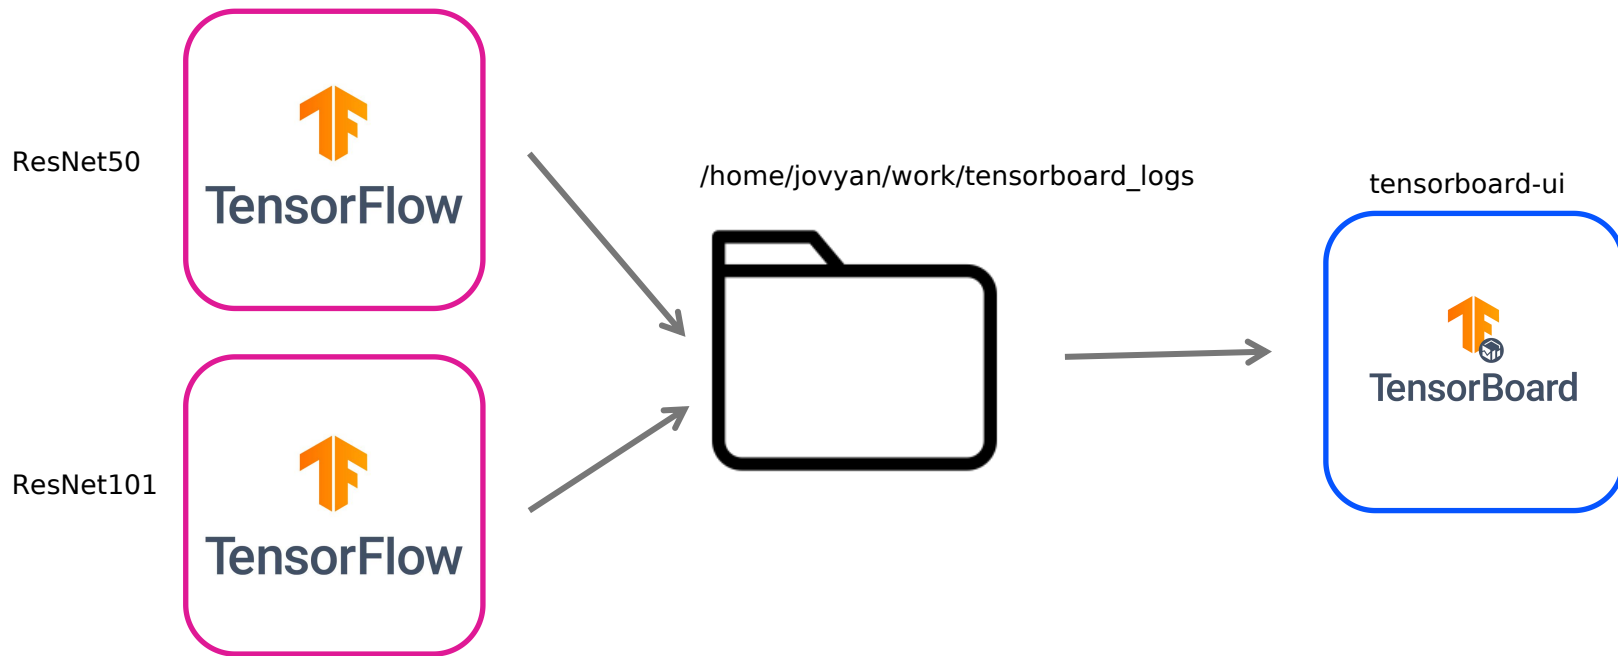
2) a new job should appear

	Job Name	Status ↓
	job-0	Running
	tb-reset-demo-0	Running

3) The Tensorboard UI will refresh periodically



Example job submission



Thank you!

